



```
In [48]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, cross_val_score, Random
from sklearn.metrics import classification_report, confusion_matrix, accuracy
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import StandardScaler
from keras.callbacks import EarlyStopping
from keras.models import Sequential
from keras.layers import Dense
import warnings
warnings.filterwarnings('ignore')
```

## Step 1 - Data Importing

```
In [49]: df = pd.read_csv(r"/Users/pvsairamsaketh/Desktop/projects/heart.csv")
df.head()
```

```
Out[49]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	c
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

```
In [50]: #checking our dependant column
df['target'].value_counts()
```

```
Out[50]: target
1      526
0      499
Name: count, dtype: int64
```

## step2 - basic data inspection

```
In [51]: df.head() #to print topp 5 rows
```

```
Out[51]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	c
<b>0</b>	52	1	0	125	212	0	1	168	0	1.0	2	
<b>1</b>	53	1	0	140	203	1	0	155	1	3.1	0	
<b>2</b>	70	1	0	145	174	0	1	125	1	2.6	0	
<b>3</b>	61	1	0	148	203	0	1	161	0	0.0	2	
<b>4</b>	62	0	0	138	294	1	1	106	0	1.9	1	

```
In [52]: df.tail()#to print bottom 5 rows
```

```
Out[52]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	c
<b>1020</b>	59	1	1	140	221	0	1	164	1	0.0		
<b>1021</b>	60	1	0	125	258	0	0	141	1	2.8		
<b>1022</b>	47	1	0	110	275	0	0	118	1	1.0		
<b>1023</b>	50	0	0	110	254	0	0	159	0	0.0		
<b>1024</b>	54	1	0	120	188	0	1	113	0	1.4		

```
In [53]: df.sample(5) #to print random 5 sample
```

```
Out[53]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	c
<b>222</b>	64	1	3	110	211	0	0	144	1	1.8	1	
<b>816</b>	70	1	1	156	245	0	0	143	0	0.0	2	
<b>649</b>	45	0	1	130	234	0	0	175	0	0.6	1	
<b>933</b>	38	1	3	120	231	0	1	182	1	3.8	1	
<b>178</b>	44	1	0	110	197	0	0	177	0	0.0	2	

```
In [54]: #to check datatypes
df.dtypes
```

```
Out[54]: age          int64
sex            int64
cp            int64
trestbps      int64
chol          int64
fbs           int64
restecg       int64
thalach       int64
exang         int64
oldpeak       float64
slope         int64
ca            int64
thal          int64
target        int64
dtype: object
```

```
In [55]: #to check information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null  int64
1   sex         1025 non-null  int64
2   cp          1025 non-null  int64
3   trestbps    1025 non-null  int64
4   chol        1025 non-null  int64
5   fbs         1025 non-null  int64
6   restecg     1025 non-null  int64
7   thalach     1025 non-null  int64
8   exang       1025 non-null  int64
9   oldpeak     1025 non-null  float64
10  slope       1025 non-null  int64
11  ca          1025 non-null  int64
12  thal        1025 non-null  int64
13  target      1025 non-null  int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
In [56]: #to check mathematical information of dataset
df.describe().T
```

Out[56]:

	count	mean	std	min	25%	50%	75%	max
<b>age</b>	1025.0	54.434146	9.072290	29.0	48.0	56.0	61.0	77.0
<b>sex</b>	1025.0	0.695610	0.460373	0.0	0.0	1.0	1.0	1.0
<b>cp</b>	1025.0	0.942439	1.029641	0.0	0.0	1.0	2.0	3.0
<b>trestbps</b>	1025.0	131.611707	17.516718	94.0	120.0	130.0	140.0	200.0
<b>chol</b>	1025.0	246.000000	51.592510	126.0	211.0	240.0	275.0	564.0
<b>fbs</b>	1025.0	0.149268	0.356527	0.0	0.0	0.0	0.0	1.0
<b>restecg</b>	1025.0	0.529756	0.527878	0.0	0.0	1.0	1.0	2.0
<b>thalach</b>	1025.0	149.114146	23.005724	71.0	132.0	152.0	166.0	202.0
<b>exang</b>	1025.0	0.336585	0.472772	0.0	0.0	0.0	1.0	1.0
<b>oldpeak</b>	1025.0	1.071512	1.175053	0.0	0.0	0.8	1.8	6.2
<b>slope</b>	1025.0	1.385366	0.617755	0.0	1.0	1.0	2.0	2.0
<b>ca</b>	1025.0	0.754146	1.030798	0.0	0.0	0.0	1.0	4.0
<b>thal</b>	1025.0	2.323902	0.620660	0.0	2.0	2.0	3.0	3.0
<b>target</b>	1025.0	0.513171	0.500070	0.0	0.0	1.0	1.0	1.0

In [57]: `df.shape`

Out[57]: (1025, 14)

## Step 3 - Data Cleaning

In [58]: `#to check null values or not`  
`df.isnull().sum() #no null values`

Out[58]:

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

```
In [59]: #Rename column names
df = df.rename(columns={'cp': 'ChestPainType', 'trestbps': 'RestingBloodPressure'})

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   1025 non-null   int64
1   sex                                   1025 non-null   int64
2   ChestPainType                         1025 non-null   int64
3   RestingBloodPressure                  1025 non-null   int64
4   Cholestrol                            1025 non-null   int64
5   FastingBloodSugar                     1025 non-null   int64
6   RestingEcg                           1025 non-null   int64
7   MaxHeartRate                         1025 non-null   int64
8   Exercise                             1025 non-null   int64
9   oldpeak                              1025 non-null   float64
10  slope                                 1025 non-null   int64
11  NoOfBloodVesselsDuringFluroscopy      1025 non-null   int64
12  Thalassemia                           1025 non-null   int64
13  target                               1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

## Step 4 - Data Analysis and Data Visualization

```
In [60]: #grouping chestpaintype , target
df1 = df.groupby('ChestPainType').agg({'target': 'count'})
```

```
In [61]: df1 #from this we can say out of 1025 records we have , 497 people got chestpa
```

```
Out[61]:
```

	target
ChestPainType	
0	497
1	167
2	284
3	77

```
In [62]: #grouping Thalassemia with Target
# Grouping Thalassemia with Target and adding percentage
df2 = (
    df.groupby(['Thalassemia', 'ChestPainType', 'NoOfBloodVesselsDuringFlurosc
    .agg(Count_of_Patients=('target', 'count'))
```

```
        .reset_index()
    )

    # Add percentage column
    df2['Percentage'] = (df2['Count_of_Patients'] / df2['Count_of_Patients'].sum())

    # Sort by Count_of_Patients descending
    df2 = df2.sort_values(by=['Count_of_Patients', 'Thalassemia', 'ChestPainType',
                             ascending=False])
```

In [63]: df2 # Most heart disease cases occur in patients with Thalassemia 2 or 3, ches

Out[63]:

	Thalassemia	ChestPainType	NoOfBloodVesselsDuringFluroscopy	Count_of_F
17	2	2		0
10	2	0		0
14	2	1		0
25	3	0		0
26	3	0		1
27	3	0		2
11	2	0		1
33	3	2		0
18	2	2		1
28	3	0		3
37	3	3		0
34	3	2		1
30	3	1		0
12	2	0		2
22	2	3		0
15	2	1		1
2	1	0		0
24	2	3		2
13	2	0		3
4	1	0		2
21	2	2		4
36	3	2		3
23	2	3		1
16	2	1		2
8	1	2		1
31	3	1		1
20	2	2		3
9	1	3		0
3	1	0		1
35	3	2		2
32	3	1		4

	Thalassemia	ChestPainType	NoOfBloodVesselsDuringFluroscopy	Count_of_F
6	1	1		0
5	1	0		3
0	0	0		0
29	3	0		4
19	2	2		2
7	1	1		3
1	0	2		0


In [64]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   1025 non-null   int64
1   sex                                   1025 non-null   int64
2   ChestPainType                        1025 non-null   int64
3   RestingBloodPressure                 1025 non-null   int64
4   Cholestrol                           1025 non-null   int64
5   FastingBloodSugar                    1025 non-null   int64
6   RestingEcg                           1025 non-null   int64
7   MaxHeartRate                         1025 non-null   int64
8   Exercise                             1025 non-null   int64
9   oldpeak                             1025 non-null   float64
10  slope                                1025 non-null   int64
11  NoOfBloodVesselsDuringFluroscopy     1025 non-null   int64
12  Thalassemia                          1025 non-null   int64
13  target                               1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```


In [65]: `import warnings`  
`warnings.filterwarnings('ignore')`

In [66]: `#to check target col count`  
`a = sns.countplot(x = 'target',data = df , palette = ['red','blue'])`  
`for ax in a.containers:`  
 `a.bar_label(ax)`

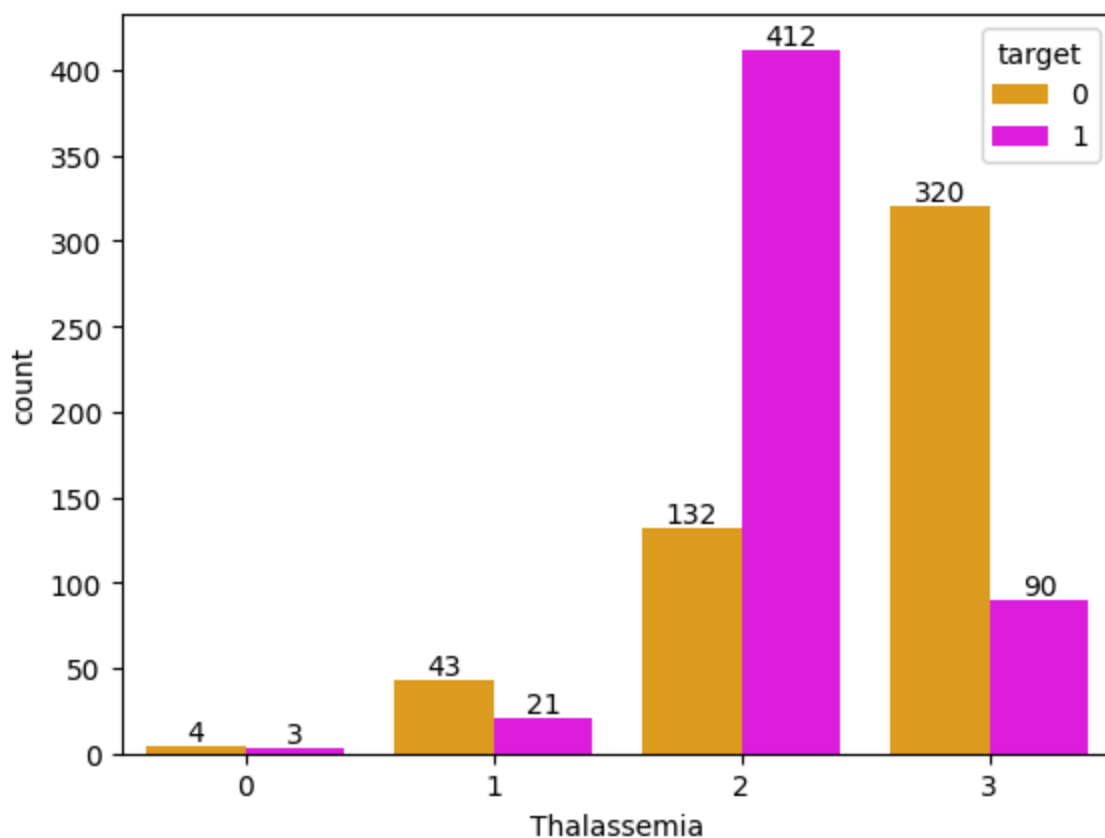




```
In [67]: #to check Thalassemia count
a = sns.countplot(x = 'Thalassemia',data = df , palette = ['green','blue','cya
for ax in a.containers:
    a.bar_label(ax)
```



```
In [68]: # to check thalassmia with target as hue
a = sns.countplot(x = 'Thalassemia',data = df ,hue = 'target',palette = ['or', 'blue', 'green', 'red'])
for ax in a.containers:
    a.bar_label(ax) #majority of them got thalassmia 2 and are prone to heart
```



```
In [69]: df.columns
```

```
Out[69]: Index(['age', 'sex', 'ChestPainType', 'RestingBloodPressure', 'Cholestrol',
               'FastingBloodSugar', 'RestingEcg', 'MaxHeartRate', 'Exercise',
               'oldpeak', 'slope', 'NoOfBloodVesselsDuringFluroscopy', 'Thalassemia',
               'target'],
              dtype='object')
```

```
In [70]: df.head()
```

```
Out[70]:
```

	age	sex	ChestPainType	RestingBloodPressure	Cholestrol	FastingBloodSugar
0	52	1	0	125	212	
1	53	1	0	140	203	
2	70	1	0	145	174	
3	61	1	0	148	203	
4	62	0	0	138	294	

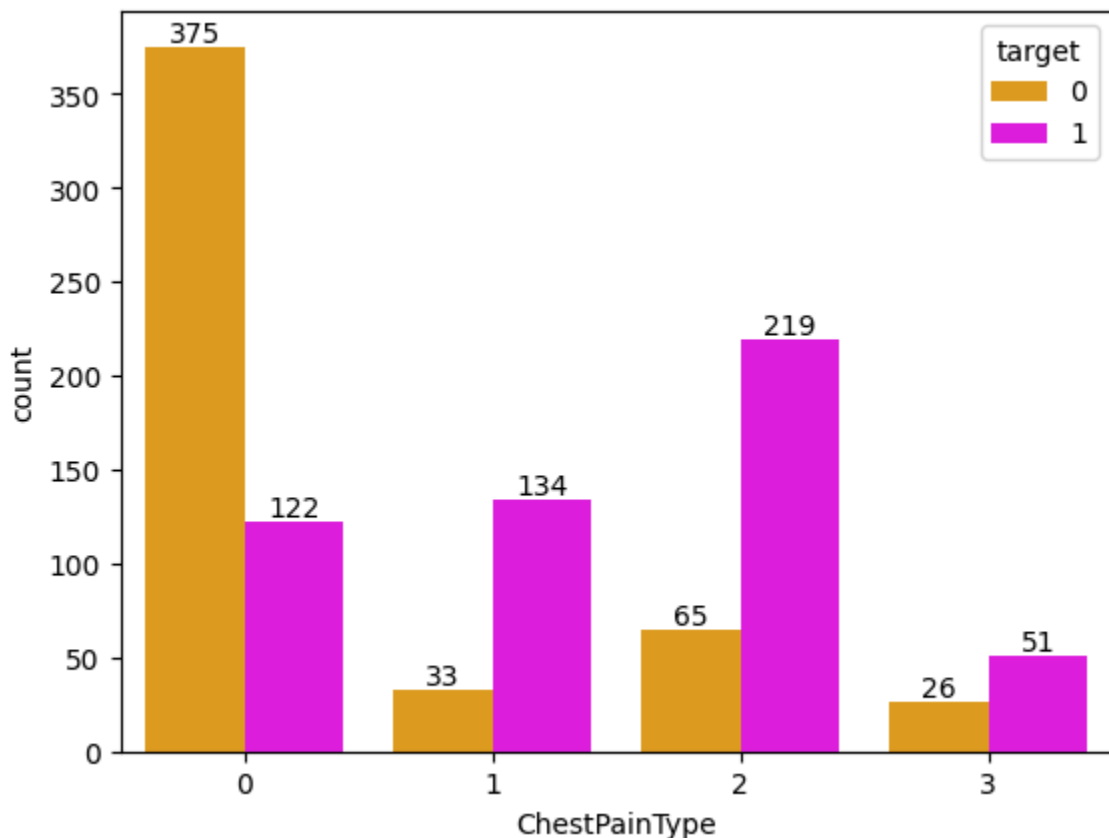
```
In [71]: df['Thalassemia'].value_counts()
```

```
Out[71]: Thalassemia
2      544
3      410
1       64
0        7
Name: count, dtype: int64
```

```
In [72]: df['NoOfBloodVesselsDuringFluroscopy'].value_counts()
```

```
Out[72]: NoOfBloodVesselsDuringFluroscopy
0      578
1      226
2      134
3       69
4       18
Name: count, dtype: int64
```

```
In [73]: #chest pain type with target
a = sns.countplot(x = 'ChestPainType',data = df ,hue = 'target',palette = ['or', 'm'])
for ax in a.containers:
    a.bar_label(ax)#people who got chesspaintype 0 , majority of them didnt go
```

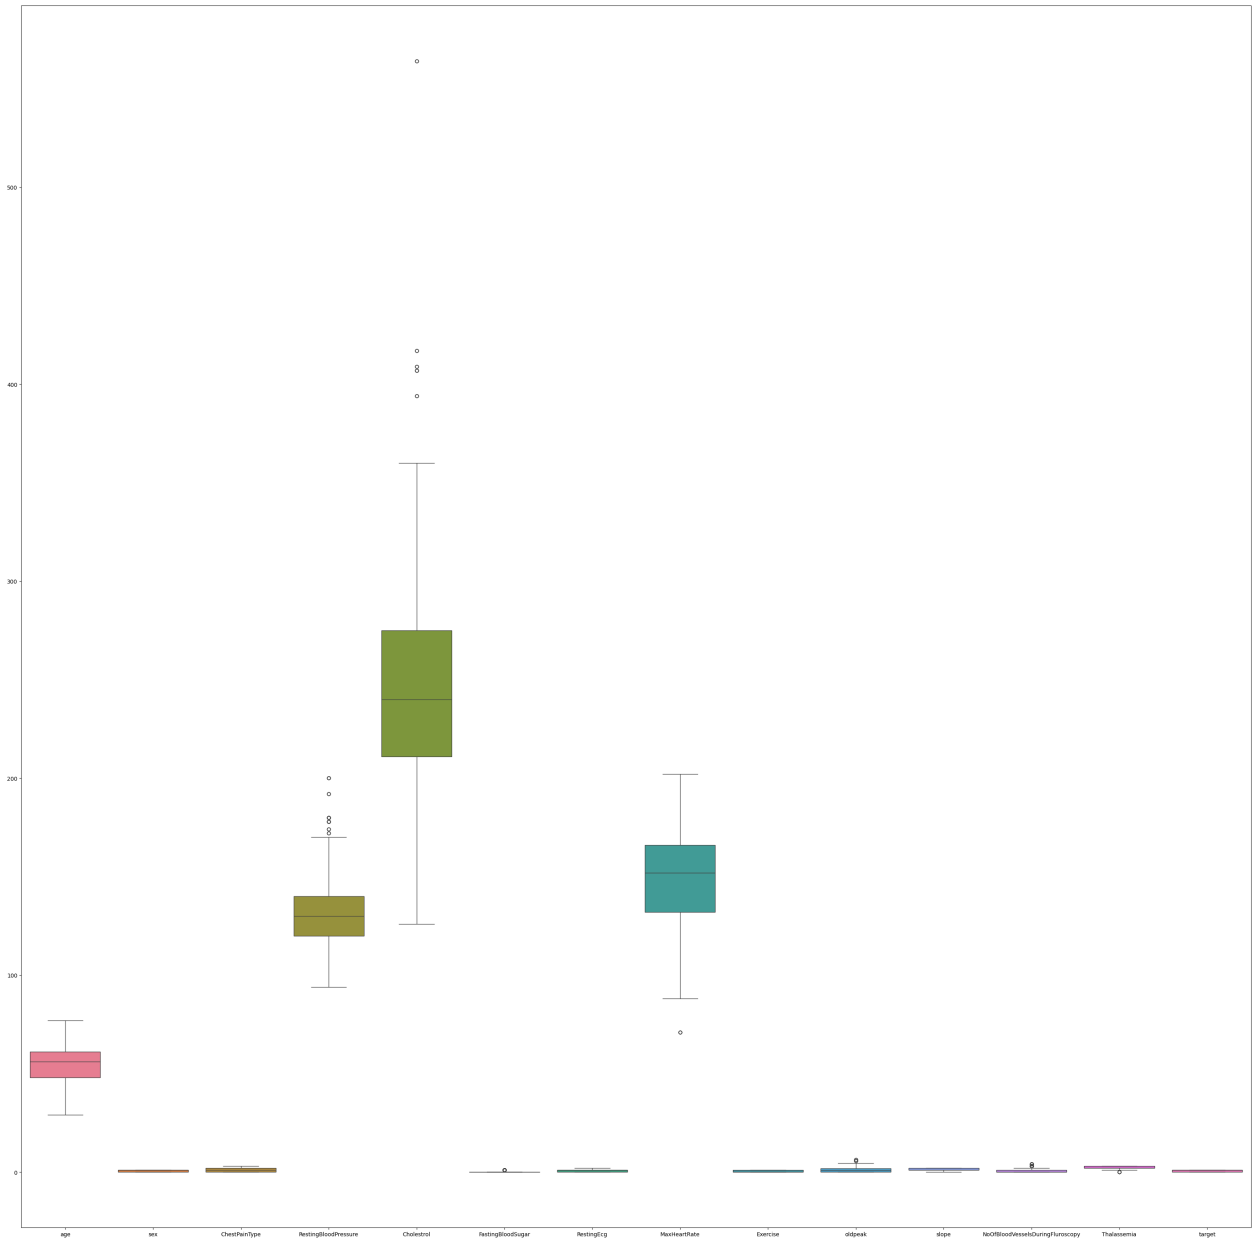


## Step 5 - Outlier Detection

```
In [74]: #to find any outliers in dataset
```

```
plt.figure(figsize=(40,40))
sns.boxplot(df)
```

Out[74]: <Axes: >



```
In [75]: #outliers are present in RestingBloodPressure , Cholestrol , FastingBloodSugar
# Columns with outliers
num_cols = ['RestingBloodPressure', 'Cholestrol', 'MaxHeartRate', 'oldpeak']
cat_ord_cols = ['FastingBloodSugar', 'NoOfBloodVesselsDuringFluoroscopy', 'Thal
```

```
In [76]: #to solve the outliers

#for numerical columns
for col in num_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
```

```

IQR = Q3 - Q1
lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR
df[col] = df[col].clip(lower, upper)
#for categorical columns

# Example: FastingBloodSugar (0 or 1)
df['FastingBloodSugar'] = df['FastingBloodSugar'].clip(0, 1)

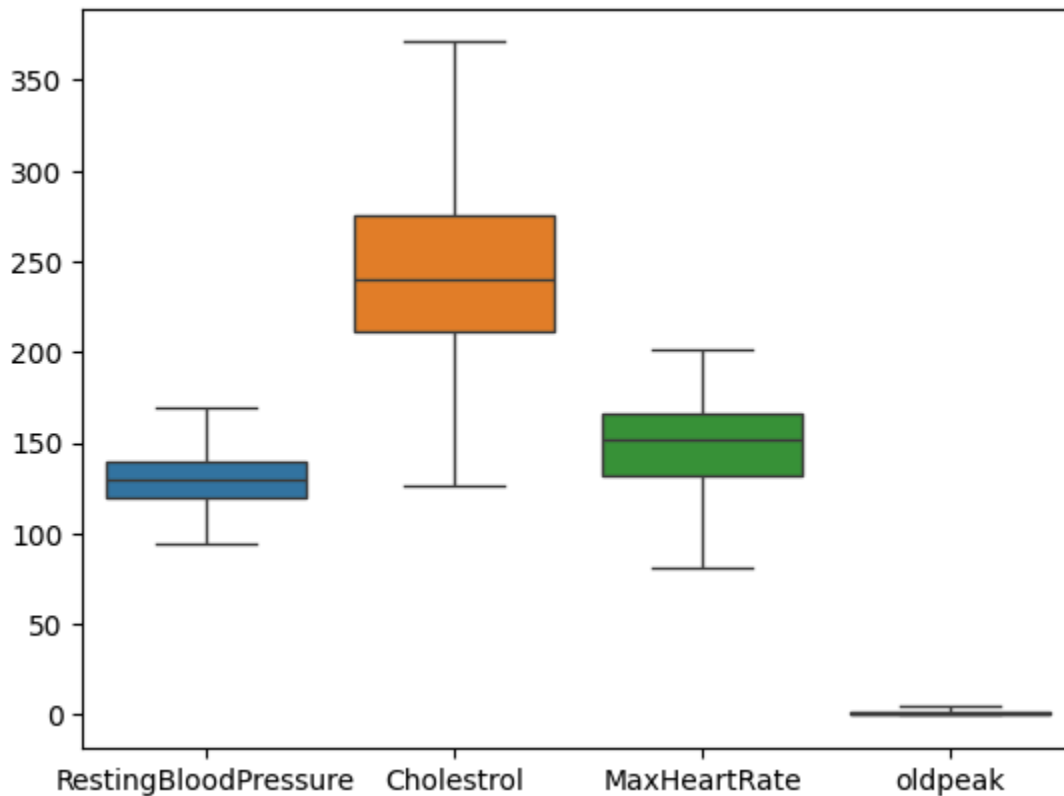
# Number of vessels (0-4)
df['NoOfBloodVesselsDuringFluroscopy'] = df['NoOfBloodVesselsDuringFluroscopy']

# Thalassemia (0-3)
df['Thalassemia'] = df['Thalassemia'].clip(0, 3)

```

```
In [77]: sns.boxplot(data=df[num_cols]) #outliers removed successfully
```

```
Out[77]: <Axes: >
```



```
In [78]: df.head()
```

```
Out[78]:
```

	age	sex	ChestPainType	RestingBloodPressure	Cholestrol	FastingBloodSugar
0	52	1	0	125	212	
1	53	1	0	140	203	
2	70	1	0	145	174	
3	61	1	0	148	203	
4	62	0	0	138	294	

```
In [79]: df['oldpeak'] = df['oldpeak'].astype('int64')
df['oldpeak'].head()
```

```
Out[79]: 0    1
1    3
2    2
3    0
4    1
Name: oldpeak, dtype: int64
```

```
In [80]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   1025 non-null   int64
1   sex                                   1025 non-null   int64
2   ChestPainType                        1025 non-null   int64
3   RestingBloodPressure                 1025 non-null   int64
4   Cholestrol                           1025 non-null   int64
5   FastingBloodSugar                   1025 non-null   int64
6   RestingEcg                           1025 non-null   int64
7   MaxHeartRate                         1025 non-null   int64
8   Exercise                             1025 non-null   int64
9   oldpeak                              1025 non-null   int64
10  slope                                1025 non-null   int64
11  NoOfBloodVesselsDuringFluroscopy     1025 non-null   int64
12  Thalassemia                           1025 non-null   int64
13  target                               1025 non-null   int64
dtypes: int64(14)
memory usage: 112.2 KB
```

## Step 6 - Model Building

```
In [81]: #splitting dataset into Dependant and Independant columns
X = df.drop('target',axis = 1)
y = df['target']
```

```

In [82]: #split dataset into 20 % test and 80% train
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size= 0.2 , random_state=42)

In [83]: len(X_train)

Out[83]: 820

In [84]: len(X_test)

Out[84]: 205

In [85]: #standard scaler
scaler = StandardScaler()
X_train= scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

In [86]: #perform models

models = [LogisticRegression(),DecisionTreeClassifier(),RandomForestClassifier()]
results = {}

In [87]: # Loop through models
for model in models:
    print(f"Model: {model.__class__.__name__}")

    # Fit model
    model.fit(X_train, y_train)

    # Predictions
    y_pred = model.predict(X_test)

    # Accuracy
    train_score = model.score(X_train, y_train)
    test_score = model.score(X_test, y_test)
    results[model.__class__.__name__] = test_score

    print(f"Training Accuracy: {train_score:.4f}")
    print(f"Test Accuracy: {test_score:.4f}")

    # Overfitting check
    if train_score > test_score + 0.05:
        print("⚠ Possible Overfitting detected!")
    elif test_score > train_score + 0.05:
        print("⚠ Possible Underfitting detected!")
    else:
        print("✅ Model seems balanced.")

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    print("Confusion Matrix:")
    print(cm)

```



```


# Optional: visualize confusion matrix
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"{model.__class__.__name__} Confusion Matrix")
plt.show()

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("-"*160)

# Optional: Compare model accuracies
plt.figure(figsize=(22,22))
#models = [LogisticRegression(),DecisionTreeClassifier(),RandomForestClassifier]
a = sns.barplot(x=list(results.keys()), y=list(results.values()),palette = ['r','b','g'])
for ax in a.containers:
    a.bar_label(ax)
plt.ylabel("Test Accuracy")
plt.title("Comparison of Model Accuracies")
plt.ylim(0,1)
plt.show()

```


Model: LogisticRegression  
 Training Accuracy: 0.8634  
 Test Accuracy: 0.7951  
 ⚠ Possible Overfitting detected!  
 Confusion Matrix:  
 [[73 29]  
 [13 90]]



Classification Report:

	precision	recall	f1-score	support
0	0.85	0.72	0.78	102
1	0.76	0.87	0.81	103
accuracy			0.80	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.80	0.79	205


-----  
-----  
--  
Model: DecisionTreeClassifier  
Training Accuracy: 1.0000  
Test Accuracy: 0.9854  
✅ Model seems balanced.  
Confusion Matrix:  
[[102 0]  
 [ 3 100]]



Classification Report:

	precision	recall	f1-score	support
0	0.97	1.00	0.99	102
1	1.00	0.97	0.99	103
accuracy			0.99	205
macro avg	0.99	0.99	0.99	205
weighted avg	0.99	0.99	0.99	205


-----  
-----  
--  
Model: RandomForestClassifier  
Training Accuracy: 1.0000  
Test Accuracy: 0.9854  
✅ Model seems balanced.  
Confusion Matrix:  
[[102 0]  
 [ 3 100]]



Classification Report:

	precision	recall	f1-score	support
0	0.97	1.00	0.99	102
1	1.00	0.97	0.99	103
accuracy			0.99	205
macro avg	0.99	0.99	0.99	205
weighted avg	0.99	0.99	0.99	205

-----  
-----  
--  
Model: SVC  
Training Accuracy: 0.9549  
Test Accuracy: 0.8878  
⚠ Possible Overfitting detected!  
Confusion Matrix:  
[[85 17]  
 [ 6 97]]



Classification Report:

	precision	recall	f1-score	support
0	0.93	0.83	0.88	102
1	0.85	0.94	0.89	103
accuracy			0.89	205
macro avg	0.89	0.89	0.89	205
weighted avg	0.89	0.89	0.89	205

-----  
-----  
--

Model: KNeighborsClassifier

Training Accuracy: 0.9439

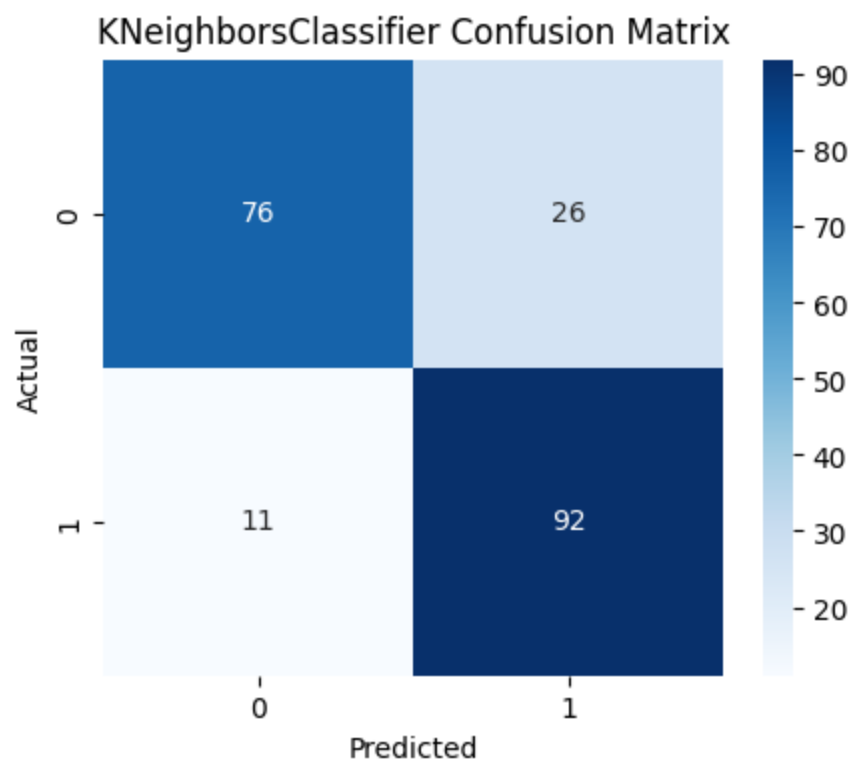
Test Accuracy: 0.8195

⚠ Possible Overfitting detected!

Confusion Matrix:

[[76 26]


[11 92]]



Classification Report:

	precision	recall	f1-score	support
0	0.87	0.75	0.80	102
1	0.78	0.89	0.83	103
accuracy			0.82	205
macro avg	0.83	0.82	0.82	205
weighted avg	0.83	0.82	0.82	205


-----  
-----  
--  
Model: GaussianNB  
Training Accuracy: 0.8512  
Test Accuracy: 0.7902  
△ Possible Overfitting detected!  
Confusion Matrix:  
[[72 30]  
 [13 90]]



Classification Report:

	precision	recall	f1-score	support
0	0.85	0.71	0.77	102
1	0.75	0.87	0.81	103
accuracy			0.79	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.79	0.79	205

-----  
-----  
--



## step 7 - Model Optimization Remove overfitting

```
In [88]: #define hyperparameters
# step 1 - Models defining
models = {
    'LogisticRegression': LogisticRegression(max_iter=10000),
    'SVC': SVC(),
    'KNeighborsClassifier': KNeighborsClassifier(),
    'GaussianNB': GaussianNB(),
    'DecisionTreeClassifier': DecisionTreeClassifier(),
    'RandomForestClassifier': RandomForestClassifier()
}
```



```

In [89]: #step 2 - defining hyperparameters
param_distributions = {
    'LogisticRegression': {
        'C': [0.1, 0.5, 1, 2],
        'penalty': ['l2', 'l1'],
        'solver': ['lbfgs', 'liblinear']
    },
    'SVC': {
        'C': [0.1, 0.5, 1, 2, 5],
        'kernel': ['linear', 'rbf'],
        'gamma': ['scale']
    },
    'KNeighborsClassifier': {
        'n_neighbors': list(range(5,16)),
        'weights': ['distance'],
        'p': [1,2]
    },
    'GaussianNB': {
        'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6]
    },
    'DecisionTreeClassifier': {
        'max_depth': list(range(3,9)),
        'min_samples_split': list(range(2,9)),
        'min_samples_leaf': list(range(1,5))
    },
    'RandomForestClassifier': {
        'n_estimators': list(range(50,151,10)),
        'max_depth': list(range(3,11)),
        'min_samples_split': list(range(2,9)),
        'min_samples_leaf': list(range(1,5))
    }
}

```

```

In [90]: # step 3 - Dictionary to store test accuracies
results = {}

```

```

In [91]: # step 4 - Hyperparameter tuning and evaluation loop
for name, model in models.items():
    print(f" ♦ Tuning {name}...")
    # step 5 - perform Randomised search and select best parameters
    rand_search = RandomizedSearchCV(estimator=model, param_distributions=param_
    rand_search.fit(X_train, y_train)
    best_model = rand_search.best_estimator_

    # Predictions
    y_pred = best_model.predict(X_test)

    # step 6 - Accuracy
    train_score = best_model.score(X_train, y_train)
    test_score = best_model.score(X_test, y_test)
    results[name] = test_score

    print(f"Best Parameters: {rand_search.best_params_}")

```

```

print(f"Training Accuracy: {train_score:.4f}")
print(f"Test Accuracy: {test_score:.4f}")

# step 7 - Overfitting check
if train_score > test_score + 0.05:
    print("⚠ Possible Overfitting detected!")
else:
    print("✅ Model seems balanced.")

# step 8 - Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"{name} Confusion Matrix")
plt.show()

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("-"*60)

# Compare model accuracies
plt.figure(figsize=(23,23))
a= sns.barplot(x=list(results.keys()), y=list(results.values()),palette = ['gr
for ax in a.containers:
    a.bar_label(ax)
plt.xlabel("Machine Learning Models")
plt.ylabel("Test Accuracy")
plt.title("Comparison of Model Accuracies After Hyperparameter Tuning")
plt.ylim(0,1)
plt.show()

```


♦ Tuning LogisticRegression...

Best Parameters: {'solver': 'lbfgs', 'penalty': 'l2', 'C': 0.1}

Training Accuracy: 0.8659

Test Accuracy: 0.8000

⚠ Possible Overfitting detected!



Classification Report:

	precision	recall	f1-score	support
0	0.86	0.72	0.78	102
1	0.76	0.88	0.82	103
accuracy			0.80	205
macro avg	0.81	0.80	0.80	205
weighted avg	0.81	0.80	0.80	205

-----


♦ Tuning SVC...

Best Parameters: {'kernel': 'rbf', 'gamma': 'scale', 'C': 5}

Training Accuracy: 0.9915

Test Accuracy: 0.9610

✅ Model seems balanced.



Classification Report:

	precision	recall	f1-score	support
0	0.97	0.95	0.96	102
1	0.95	0.97	0.96	103
accuracy			0.96	205
macro avg	0.96	0.96	0.96	205
weighted avg	0.96	0.96	0.96	205


-----  
♦ Tuning KNeighborsClassifier...

Best Parameters: {'weights': 'distance', 'p': 1, 'n\_neighbors': 10}

Training Accuracy: 1.0000

Test Accuracy: 1.0000


✅ Model seems balanced.



Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	102
1	1.00	1.00	1.00	103
accuracy			1.00	205
macro avg	1.00	1.00	1.00	205
weighted avg	1.00	1.00	1.00	205

-----  
♦ Tuning GaussianNB...  
Best Parameters: {'var\_smoothing': 1e-09}  
Training Accuracy: 0.8512  
Test Accuracy: 0.7902  
⚠ Possible Overfitting detected!



Classification Report:

	precision	recall	f1-score	support
0	0.85	0.71	0.77	102
1	0.75	0.87	0.81	103
accuracy			0.79	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.79	0.79	205


-----  
♦ Tuning DecisionTreeClassifier...

Best Parameters: {'min\_samples\_split': 4, 'min\_samples\_leaf': 1, 'max\_depth': 8}

Training Accuracy: 0.9927

Test Accuracy: 0.9610

✅ Model seems balanced.



Classification Report:

	precision	recall	f1-score	support
0	0.97	0.95	0.96	102
1	0.95	0.97	0.96	103
accuracy			0.96	205
macro avg	0.96	0.96	0.96	205
weighted avg	0.96	0.96	0.96	205

-----


♦ Tuning RandomForestClassifier...

Best Parameters: {'n\_estimators': 130, 'min\_samples\_split': 3, 'min\_samples\_leaf': 1, 'max\_depth': 10}

Training Accuracy: 1.0000

Test Accuracy: 0.9854

✅ Model seems balanced.




Classification Report:

	precision	recall	f1-score	support
0	0.97	1.00	0.99	102
1	1.00	0.97	0.99	103
accuracy			0.99	205
macro avg	0.99	0.99	0.99	205
weighted avg	0.99	0.99	0.99	205

-----





In [92]: *#knn is the best method - test it*

```
# Best KNN model
best_knn = KNeighborsClassifier(weights='distance', p=1, n_neighbors=10)
best_knn.fit(X_train, y_train)
y_pred = best_knn.predict(X_test)

print("Test Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Test Accuracy: 1.0

Confusion Matrix:

```
[[102  0]
```

```
[ 0 103]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	102
1	1.00	1.00	1.00	103
accuracy			1.00	205
macro avg	1.00	1.00	1.00	205
weighted avg	1.00	1.00	1.00	205

## step 8 - Deep Learning model

### metrics types

Metric	Type	Description / What It Measures	Use Case / Example
<b>accuracy</b>	Classification	Overall fraction of correctly predicted samples.	General classification tasks.
<b>binary_accuracy</b>	Classification	Accuracy for binary classification (0 or 1).	Binary output (e.g., disease yes/no).
<b>categorical_accuracy</b>	Classification	Accuracy for multi-class with one-hot labels.	Multi-class (e.g., image classification).
<b>sparse_categorical_accuracy</b>	Classification	Accuracy for multi-class with integer labels.	Multi-class (without one-hot encoding).
<b>Precision</b>	Classification	% of correct positive predictions among all predicted positives.	When false positives are costly (e.g., spam detection).
<b>Recall</b>	Classification	% of actual positives correctly predicted.	When missing positives is costly (e.g., cancer detection).

Metric	Type	Description / What It Measures	Use Case / Example
<b>F1-Score</b>	Classification	Harmonic mean of precision and recall.	Balanced measure for imbalanced data.
<b>AUC (Area Under Curve)</b>	Classification	Measures ability to distinguish between classes.	Evaluating model performance (ROC curve).
<b>TruePositives</b>	Classification	Number of correctly identified positive samples.	Model diagnostics.
<b>TrueNegatives</b>	Classification	Number of correctly identified negatives.	Model diagnostics.
<b>FalsePositives</b>	Classification	Incorrectly predicted positives.	Helps understand model errors.
<b>FalseNegatives</b>	Classification	Missed positive cases.	Important in medical or fraud detection.
<b>MeanSquaredError (MSE)</b>	Regression	Average squared difference between actual and predicted.	Continuous outputs (e.g., price prediction).
<b>MeanAbsoluteError (MAE)</b>	Regression	Average absolute difference between actual and predicted.	Regression — less sensitive to outliers.
<b>RootMeanSquaredError (RMSE)</b>	Regression	Square root of MSE; penalizes large errors more.	Regression with emphasis on large errors.
<b>R<sup>2</sup> (Coefficient of Determination)</b>	Regression	Measures proportion of variance explained by model.	Regression performance measure.
<b>MeanAbsolutePercentageError (MAPE)</b>	Regression	Average % difference between prediction and actual value.	Forecasting models.
<b>cosine_similarity</b>	Similarity	Measures similarity between two	NLP, embeddings, recommendation

Metric	Type	Description / What It Measures	Use Case / Example
		vectors (ranges -1 to 1).	systems.
<code>log_cosh_error</code>	Regression	Smooth alternative to MSE; robust to outliers.	Regression tasks requiring smooth gradients.

## optimizer types -

Optimizer	Description
<code>SGD</code>	Stochastic Gradient Descent — updates weights using gradients; simple but may converge slowly.
<code>SGD (with Momentum)</code>	Adds momentum to accelerate learning and avoid local minima.
<code>Nesterov</code>	A variant of momentum that looks ahead before updating parameters for faster convergence.
<code>Adagrad</code>	Adapts learning rates for each parameter; good for sparse data but learning rate decays quickly.
<code>RMSprop</code>	Adapts learning rate using a moving average of squared gradients; good for non-stationary data.
<code>Adam</code>	Combines Momentum and RMSprop; most widely used optimizer in deep learning.
<code>Adamax</code>	Variant of Adam using the infinity norm for better stability in some cases.
<code>Nadam</code>	Adam + Nesterov Momentum; slightly faster convergence than Adam.
<code>Ftrl</code>	Useful for large-scale linear models; supports L1 and L2 regularization.
<code>Adadelat</code>	Extension of Adagrad that reduces its aggressive, monotonically decreasing learning rate.
<code>Lion</code>	New optimizer that uses sign-based gradient updates for faster and more stable training.

# Loss types

Loss Function	Type	Description	When to Use / Example
<code>binary_crossentropy</code>	Classification	Measures the difference between two probability distributions for binary (0/1) classification.	Binary classification (e.g., heart disease prediction).
<code>categorical_crossentropy</code>	Classification	Used for multi-class classification with one-hot encoded labels.	Multi-class classification (e.g., digit recognition with one-hot labels).
<code>sparse_categorical_crossentropy</code>	Classification	Same as above but for integer class labels (not one-hot).	Multi-class classification with integer targets.
<code>mean_squared_error</code> (MSE)	Regression	Measures average squared difference between actual and predicted values.	Regression tasks (e.g., predicting prices, continuous outputs).
<code>mean_absolute_error</code> (MAE)	Regression	Measures average absolute difference between actual and predicted values.	Regression, less sensitive to outliers than MSE.
<code>hinge</code>	Classification	Used for “maximum-margin” classification (like SVMs).	Binary classification with $\{-1, +1\}$ labels.
<code>squared_hinge</code>	Classification	Square of hinge loss; penalizes larger margin	Similar to hinge, but smoother gradients.

Loss Function	Type	Description	When to Use / Example
		violations more strongly.	
kullback_leibler_divergence (KL Divergence)	Probabilistic Models	Measures how one probability distribution diverges from another.	Variational autoencoders (VAEs), probabilistic models.
poisson	Regression / Count Data	Suitable for modeling count data (e.g., event counts).	When output represents count rates (e.g., number of calls, clicks).
cosine_similarity	Similarity-Based	Measures cosine similarity between predicted and true values.	NLP, embeddings, or recommendation systems.
huber	Regression	Combination of MSE and MAE — less sensitive to outliers.	Robust regression with outlier resistance.
log_cosh	Regression	Logarithm of hyperbolic cosine of prediction error; smooth alternative to MSE.	When you want smooth gradients and robustness.

| categorical\_hinge | Classification | Hinge loss for multi-class problems. | Multi-class problems with {-1, +1} encoded labels.

# Activation Functions

Activation Function	Formula	Range	Description	When to Use
<b>Sigmoid (<math>\sigma</math>)</b>	$\sigma(x) = 1 / (1 + e^{(-x)})$	(0, 1)	Converts any real value into a probability-like output.	Output layer for <b>binary classification</b> .
<b>Tanh</b>	$\tanh(x) = (e^x - e^{-x}) / (e^x + e^{-x})$	(-1, 1)	Zero-centered version	Hidden layers

Activation Function	Formula	Range	Description	When to Use
<b>(Hyperbolic Tangent)</b>	$-e^{(-x)} / (e^{(x)} + e^{(-x)})$		of Sigmoid; stronger gradients.	when inputs are <b>normalized</b> .
<b>ReLU (Rectified Linear Unit)</b>	$\text{ReLU}(x) = \max(0, x)$	$[0, \infty)$	Fast and efficient; avoids vanishing gradient problem.	Default choice for <b>hidden layers</b> .
<b>Leaky ReLU</b>	$f(x) = x \text{ if } x > 0 \text{ else } 0.01x$	$(-\infty, \infty)$	Solves “dead neuron” issue by allowing a small negative slope.	When some neurons die with ReLU.
<b>Parametric ReLU (PReLU)</b>	$f(x) = x \text{ if } x > 0 \text{ else } \alpha x$	$(-\infty, \infty)$	Learns the slope ( $\alpha$ ) automatically for better performance.	Used in <b>deep CNNs</b> .
<b>ELU (Exponential Linear Unit)</b>	$f(x) = x \text{ if } x > 0 \text{ else } \alpha*(e^x - 1)$	$(-\alpha, \infty)$	Helps achieve faster and more stable learning.	For <b>deep networks</b> .
<b>SELU (Scaled ELU)</b>	$f(x) = \lambda * (x \text{ if } x > 0 \text{ else } \alpha*(e^x - 1))$	$(-\infty, \infty)$	Self-normalizing activation, stabilizes training.	When using <b>Self-Normalizing Neural Nets</b> .
<b>Softmax</b>	$f(x_i) = e^{(x_i)} / \sum e^{(x_j)}$	$(0, 1)$	Converts logits to probabilities across multiple classes.	Output layer for <b>multi-class classification</b> .
<b>Swish</b>	$f(x) = x * \text{sigmoid}(x)$	$(-\infty, \infty)$	Smooth and non-monotonic; better than ReLU in deep models.	For <b>deep neural networks</b> .
<b>GELU (Gaussian Error Linear Unit)</b>	$f(x) = x * \Phi(x)$	$(-\infty, \infty)$	Used in Transformers; smooth like Swish.	<b>NLP models</b> (e.g., BERT, GPT).
<b>Softplus</b>	$f(x) = \ln(1 + e^x)$	$(0, \infty)$	Smooth approximation of ReLU; differentiable everywhere.	When smooth output is required.
<b>Hard Sigmoid</b>	$f(x) = \text{clip}((x * 0.2) + 0.5, 0, 1)$	$(0, 1)$	Fast, piecewise linear version of Sigmoid.	For <b>mobile/low-power models</b> .
<b>Hard Swish</b>	$f(x) = x * \text{ReLU6}(x + 3) / 6$	$(-\infty, \infty)$	Lightweight and efficient; used in MobileNetV3.	For <b>mobile and embedded models</b> .

```
In [93]: # Define ANN model
# =====
model = Sequential([
    Dense(64, input_dim=X_train.shape[1], activation='relu'), # Hidden layer
    Dense(32, activation='relu'), # Hidden layer
    Dense(16, activation='relu'), # Hidden layer
    Dense(1, activation='sigmoid') # Output layer
])
```

```
In [94]: # Compile the model
# Compile model
# =====
model.compile(optimizer="adam", loss='binary_crossentropy', metrics=['accuracy
```

```
In [95]: # Train the model
# =====
history = model.fit(
    X_train, y_train,
    validation_split=0.2,
    epochs=100,
    batch_size=16,
    verbose=2 # One line per epoch
)
```



Epoch 1/100  
41/41 - 1s - 21ms/step - accuracy: 0.7820 - binary\_accuracy: 0.7820 - loss: 0.5543 - recall: 0.7626 - val\_accuracy: 0.8171 - val\_binary\_accuracy: 0.8171 - val\_loss: 0.4691 - val\_recall: 0.8023

Epoch 2/100  
41/41 - 0s - 1ms/step - accuracy: 0.8704 - binary\_accuracy: 0.8704 - loss: 0.3601 - recall: 0.9050 - val\_accuracy: 0.8232 - val\_binary\_accuracy: 0.8232 - val\_loss: 0.4131 - val\_recall: 0.8721

Epoch 3/100  
41/41 - 0s - 1ms/step - accuracy: 0.8918 - binary\_accuracy: 0.8918 - loss: 0.2867 - recall: 0.9318 - val\_accuracy: 0.8110 - val\_binary\_accuracy: 0.8110 - val\_loss: 0.4079 - val\_recall: 0.8721

Epoch 4/100  
41/41 - 0s - 1ms/step - accuracy: 0.9040 - binary\_accuracy: 0.9040 - loss: 0.2531 - recall: 0.9318 - val\_accuracy: 0.8659 - val\_binary\_accuracy: 0.8659 - val\_loss: 0.3777 - val\_recall: 0.8837

Epoch 5/100  
41/41 - 0s - 971us/step - accuracy: 0.9177 - binary\_accuracy: 0.9177 - loss: 0.2324 - recall: 0.9318 - val\_accuracy: 0.8232 - val\_binary\_accuracy: 0.8232 - val\_loss: 0.3945 - val\_recall: 0.9070

Epoch 6/100  
41/41 - 0s - 927us/step - accuracy: 0.9253 - binary\_accuracy: 0.9253 - loss: 0.2087 - recall: 0.9496 - val\_accuracy: 0.8659 - val\_binary\_accuracy: 0.8659 - val\_loss: 0.3563 - val\_recall: 0.8837

Epoch 7/100  
41/41 - 0s - 929us/step - accuracy: 0.9268 - binary\_accuracy: 0.9268 - loss: 0.1981 - recall: 0.9407 - val\_accuracy: 0.8780 - val\_binary\_accuracy: 0.8780 - val\_loss: 0.3499 - val\_recall: 0.9070

Epoch 8/100  
41/41 - 0s - 928us/step - accuracy: 0.9375 - binary\_accuracy: 0.9375 - loss: 0.1797 - recall: 0.9525 - val\_accuracy: 0.8841 - val\_binary\_accuracy: 0.8841 - val\_loss: 0.3274 - val\_recall: 0.9070

Epoch 9/100  
41/41 - 0s - 959us/step - accuracy: 0.9466 - binary\_accuracy: 0.9466 - loss: 0.1581 - recall: 0.9555 - val\_accuracy: 0.8902 - val\_binary\_accuracy: 0.8902 - val\_loss: 0.3270 - val\_recall: 0.9070

Epoch 10/100  
41/41 - 0s - 996us/step - accuracy: 0.9512 - binary\_accuracy: 0.9512 - loss: 0.1451 - recall: 0.9525 - val\_accuracy: 0.8841 - val\_binary\_accuracy: 0.8841 - val\_loss: 0.3191 - val\_recall: 0.8953

Epoch 11/100  
41/41 - 0s - 971us/step - accuracy: 0.9527 - binary\_accuracy: 0.9527 - loss: 0.1296 - recall: 0.9585 - val\_accuracy: 0.8902 - val\_binary\_accuracy: 0.8902 - val\_loss: 0.3182 - val\_recall: 0.9070

Epoch 12/100  
41/41 - 0s - 960us/step - accuracy: 0.9588 - binary\_accuracy: 0.9588 - loss: 0.1174 - recall: 0.9585 - val\_accuracy: 0.8963 - val\_binary\_accuracy: 0.8963 - val\_loss: 0.3077 - val\_recall: 0.9186

Epoch 13/100  
41/41 - 0s - 947us/step - accuracy: 0.9726 - binary\_accuracy: 0.9726 - loss: 0.1059 - recall: 0.9822 - val\_accuracy: 0.8963 - val\_binary\_accuracy: 0.8963 - val\_loss: 0.2884 - val\_recall: 0.9186

Epoch 14/100  
41/41 - 0s - 982us/step - accuracy: 0.9771 - binary\_accuracy: 0.9771 - loss:

0.0930 - recall: 0.9852 - val\_accuracy: 0.8902 - val\_binary\_accuracy: 0.8902 -  
val\_loss: 0.2801 - val\_recall: 0.8837  
Epoch 15/100  
41/41 - 0s - 983us/step - accuracy: 0.9802 - binary\_accuracy: 0.9802 - loss:  
0.0841 - recall: 0.9881 - val\_accuracy: 0.9024 - val\_binary\_accuracy: 0.9024 -  
val\_loss: 0.2795 - val\_recall: 0.9302  
Epoch 16/100  
41/41 - 0s - 977us/step - accuracy: 0.9802 - binary\_accuracy: 0.9802 - loss:  
0.0744 - recall: 0.9881 - val\_accuracy: 0.9024 - val\_binary\_accuracy: 0.9024 -  
val\_loss: 0.2856 - val\_recall: 0.9302  
Epoch 17/100  
41/41 - 0s - 955us/step - accuracy: 0.9832 - binary\_accuracy: 0.9832 - loss:  
0.0651 - recall: 0.9941 - val\_accuracy: 0.9085 - val\_binary\_accuracy: 0.9085 -  
val\_loss: 0.2629 - val\_recall: 0.9419  
Epoch 18/100  
41/41 - 0s - 952us/step - accuracy: 0.9909 - binary\_accuracy: 0.9909 - loss:  
0.0538 - recall: 0.9941 - val\_accuracy: 0.9146 - val\_binary\_accuracy: 0.9146 -  
val\_loss: 0.2465 - val\_recall: 0.9070  
Epoch 19/100  
41/41 - 0s - 935us/step - accuracy: 0.9939 - binary\_accuracy: 0.9939 - loss:  
0.0478 - recall: 0.9941 - val\_accuracy: 0.9329 - val\_binary\_accuracy: 0.9329 -  
val\_loss: 0.2496 - val\_recall: 0.9419  
Epoch 20/100  
41/41 - 0s - 955us/step - accuracy: 0.9924 - binary\_accuracy: 0.9924 - loss:  
0.0413 - recall: 0.9941 - val\_accuracy: 0.9329 - val\_binary\_accuracy: 0.9329 -  
val\_loss: 0.2542 - val\_recall: 0.9419  
Epoch 21/100  
41/41 - 0s - 990us/step - accuracy: 0.9939 - binary\_accuracy: 0.9939 - loss:  
0.0390 - recall: 0.9941 - val\_accuracy: 0.9329 - val\_binary\_accuracy: 0.9329 -  
val\_loss: 0.2394 - val\_recall: 0.9419  
Epoch 22/100  
41/41 - 0s - 1ms/step - accuracy: 0.9939 - binary\_accuracy: 0.9939 - loss: 0.03  
26 - recall: 0.9941 - val\_accuracy: 0.9268 - val\_binary\_accuracy: 0.9268 - va  
l\_loss: 0.2372 - val\_recall: 0.9070  
Epoch 23/100  
41/41 - 0s - 1ms/step - accuracy: 0.9939 - binary\_accuracy: 0.9939 - loss: 0.02  
84 - recall: 0.9941 - val\_accuracy: 0.9451 - val\_binary\_accuracy: 0.9451 - va  
l\_loss: 0.2380 - val\_recall: 0.9419  
Epoch 24/100  
41/41 - 0s - 1ms/step - accuracy: 0.9954 - binary\_accuracy: 0.9954 - loss: 0.02  
60 - recall: 0.9970 - val\_accuracy: 0.9451 - val\_binary\_accuracy: 0.9451 - va  
l\_loss: 0.2221 - val\_recall: 0.9419  
Epoch 25/100  
41/41 - 0s - 1ms/step - accuracy: 0.9970 - binary\_accuracy: 0.9970 - loss: 0.02  
15 - recall: 1.0000 - val\_accuracy: 0.9329 - val\_binary\_accuracy: 0.9329 - va  
l\_loss: 0.2247 - val\_recall: 0.9419  
Epoch 26/100  
41/41 - 0s - 1ms/step - accuracy: 0.9970 - binary\_accuracy: 0.9970 - loss: 0.02  
00 - recall: 1.0000 - val\_accuracy: 0.9329 - val\_binary\_accuracy: 0.9329 - va  
l\_loss: 0.2317 - val\_recall: 0.9419  
Epoch 27/100  
41/41 - 0s - 1ms/step - accuracy: 0.9970 - binary\_accuracy: 0.9970 - loss: 0.01  
66 - recall: 1.0000 - val\_accuracy: 0.9451 - val\_binary\_accuracy: 0.9451 - va  
l\_loss: 0.2270 - val\_recall: 0.9419

Epoch 28/100  
41/41 - 0s - 947us/step - accuracy: 0.9985 - binary\_accuracy: 0.9985 - loss: 0.0132 - recall: 1.0000 - val\_accuracy: 0.9451 - val\_binary\_accuracy: 0.9451 - val\_loss: 0.2284 - val\_recall: 0.9419

Epoch 29/100  
41/41 - 0s - 973us/step - accuracy: 0.9985 - binary\_accuracy: 0.9985 - loss: 0.0127 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2332 - val\_recall: 0.9419

Epoch 30/100  
41/41 - 0s - 996us/step - accuracy: 0.9985 - binary\_accuracy: 0.9985 - loss: 0.0111 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2268 - val\_recall: 0.9419

Epoch 31/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0097 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2308 - val\_recall: 0.9419

Epoch 32/100  
41/41 - 0s - 991us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0083 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2410 - val\_recall: 0.9419

Epoch 33/100  
41/41 - 0s - 970us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0072 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2376 - val\_recall: 0.9419

Epoch 34/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0061 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2432 - val\_recall: 0.9419

Epoch 35/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0061 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2446 - val\_recall: 0.9419

Epoch 36/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0049 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2463 - val\_recall: 0.9419

Epoch 37/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0044 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2500 - val\_recall: 0.9419

Epoch 38/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0044 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2485 - val\_recall: 0.9419

Epoch 39/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0035 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2565 - val\_recall: 0.9419

Epoch 40/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0031 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2617 - val\_recall: 0.9419

Epoch 41/100  
41/41 - 0s - 993us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:

0.0028 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2575 - val\_recall: 0.9419  
Epoch 42/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
26 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2637 - val\_recall: 0.9419  
Epoch 43/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
24 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2686 - val\_recall: 0.9419  
Epoch 44/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
22 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2701 - val\_recall: 0.9419  
Epoch 45/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
19 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2682 - val\_recall: 0.9419  
Epoch 46/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
19 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2731 - val\_recall: 0.9419  
Epoch 47/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
17 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2726 - val\_recall: 0.9419  
Epoch 48/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.00  
17 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - va  
l\_loss: 0.2778 - val\_recall: 0.9419  
Epoch 49/100  
41/41 - 0s - 971us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:  
0.0015 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2796 - val\_recall: 0.9419  
Epoch 50/100  
41/41 - 0s - 949us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:  
0.0014 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2798 - val\_recall: 0.9419  
Epoch 51/100  
41/41 - 0s - 967us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:  
0.0013 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2818 - val\_recall: 0.9419  
Epoch 52/100  
41/41 - 0s - 941us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:  
0.0012 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2826 - val\_recall: 0.9419  
Epoch 53/100  
41/41 - 0s - 958us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:  
0.0012 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2865 - val\_recall: 0.9419  
Epoch 54/100  
41/41 - 0s - 918us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:  
0.0011 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 -  
val\_loss: 0.2855 - val\_recall: 0.9419

Epoch 55/100  
41/41 - 0s - 932us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 0.0010 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2900 - val\_recall: 0.9419

Epoch 56/100  
41/41 - 0s - 920us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 9.3919e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2921 - val\_recall: 0.9419

Epoch 57/100  
41/41 - 0s - 925us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 9.1104e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2953 - val\_recall: 0.9419

Epoch 58/100  
41/41 - 0s - 935us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 8.2235e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2945 - val\_recall: 0.9419

Epoch 59/100  
41/41 - 0s - 949us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 7.8962e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2991 - val\_recall: 0.9419

Epoch 60/100  
41/41 - 0s - 936us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 7.5271e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.2998 - val\_recall: 0.9419

Epoch 61/100  
41/41 - 0s - 930us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 6.9982e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3002 - val\_recall: 0.9419

Epoch 62/100  
41/41 - 0s - 946us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 6.8541e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3009 - val\_recall: 0.9419

Epoch 63/100  
41/41 - 0s - 946us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 6.3750e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3064 - val\_recall: 0.9419

Epoch 64/100  
41/41 - 0s - 949us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 5.8996e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3069 - val\_recall: 0.9419

Epoch 65/100  
41/41 - 0s - 960us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 5.7487e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3097 - val\_recall: 0.9419

Epoch 66/100  
41/41 - 0s - 934us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 5.5162e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3100 - val\_recall: 0.9419

Epoch 67/100  
41/41 - 0s - 930us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 5.0438e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3123 - val\_recall: 0.9419

Epoch 68/100  
41/41 - 0s - 944us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss:

4.8593e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3136 - val\_recall: 0.9419  
Epoch 69/100  
41/41 - 0s - 940us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 4.6734e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3128 - val\_recall: 0.9419  
Epoch 70/100  
41/41 - 0s - 966us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 4.5426e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3164 - val\_recall: 0.9419  
Epoch 71/100  
41/41 - 0s - 940us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 4.2147e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3182 - val\_recall: 0.9419  
Epoch 72/100  
41/41 - 0s - 954us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 4.0427e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3206 - val\_recall: 0.9419  
Epoch 73/100  
41/41 - 0s - 943us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.8521e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3208 - val\_recall: 0.9419  
Epoch 74/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.635e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3230 - val\_recall: 0.9419  
Epoch 75/100  
41/41 - 0s - 967us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.5093e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3264 - val\_recall: 0.9419  
Epoch 76/100  
41/41 - 0s - 970us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.3336e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3260 - val\_recall: 0.9419  
Epoch 77/100  
41/41 - 0s - 949us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.2414e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3283 - val\_recall: 0.9419  
Epoch 78/100  
41/41 - 0s - 955us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.1382e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3311 - val\_recall: 0.9419  
Epoch 79/100  
41/41 - 0s - 954us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 3.0036e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3300 - val\_recall: 0.9419  
Epoch 80/100  
41/41 - 0s - 955us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.9856e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3343 - val\_recall: 0.9419  
Epoch 81/100  
41/41 - 0s - 968us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.7300e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3333 - val\_recall: 0.9419

Epoch 82/100  
41/41 - 0s - 968us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.5860e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3348 - val\_recall: 0.9419

Epoch 83/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.21e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3361 - val\_recall: 0.9419

Epoch 84/100  
41/41 - 0s - 972us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.4172e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3374 - val\_recall: 0.9419

Epoch 85/100  
41/41 - 0s - 976us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.4022e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3377 - val\_recall: 0.9419

Epoch 86/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.2060e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3407 - val\_recall: 0.9419

Epoch 87/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.0920e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3417 - val\_recall: 0.9419

Epoch 88/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 2.0810e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3436 - val\_recall: 0.9419

Epoch 89/100  
41/41 - 0s - 992us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.9952e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3440 - val\_recall: 0.9419

Epoch 90/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.9734e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3477 - val\_recall: 0.9419

Epoch 91/100  
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.8383e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3458 - val\_recall: 0.9419

Epoch 92/100  
41/41 - 0s - 957us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.7834e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3485 - val\_recall: 0.9419

Epoch 93/100  
41/41 - 0s - 927us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.6843e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3494 - val\_recall: 0.9419

Epoch 94/100  
41/41 - 0s - 902us/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.5997e-04 - recall: 1.0000 - val\_accuracy: 0.9573 - val\_binary\_accuracy: 0.9573 - val\_loss: 0.3509 - val\_recall: 0.9419

Epoch 95/100  
41/41 - 0s - 2ms/step - accuracy: 1.0000 - binary\_accuracy: 1.0000 - loss: 1.56

```

15e-04 - recall: 1.0000 - val_accuracy: 0.9573 - val_binary_accuracy: 0.9573 -
val_loss: 0.3516 - val_recall: 0.9419
Epoch 96/100
41/41 - 0s - 990us/step - accuracy: 1.0000 - binary_accuracy: 1.0000 - loss:
1.5042e-04 - recall: 1.0000 - val_accuracy: 0.9573 - val_binary_accuracy: 0.957
3 - val_loss: 0.3527 - val_recall: 0.9419
Epoch 97/100
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary_accuracy: 1.0000 - loss: 1.47
74e-04 - recall: 1.0000 - val_accuracy: 0.9573 - val_binary_accuracy: 0.9573 -
val_loss: 0.3552 - val_recall: 0.9419
Epoch 98/100
41/41 - 0s - 1ms/step - accuracy: 1.0000 - binary_accuracy: 1.0000 - loss: 1.39
88e-04 - recall: 1.0000 - val_accuracy: 0.9573 - val_binary_accuracy: 0.9573 -
val_loss: 0.3554 - val_recall: 0.9419
Epoch 99/100
41/41 - 0s - 987us/step - accuracy: 1.0000 - binary_accuracy: 1.0000 - loss:
1.3613e-04 - recall: 1.0000 - val_accuracy: 0.9573 - val_binary_accuracy: 0.957
3 - val_loss: 0.3561 - val_recall: 0.9419
Epoch 100/100
41/41 - 0s - 982us/step - accuracy: 1.0000 - binary_accuracy: 1.0000 - loss:
1.3084e-04 - recall: 1.0000 - val_accuracy: 0.9573 - val_binary_accuracy: 0.957
3 - val_loss: 0.3584 - val_recall: 0.9419

```

```

In [96]: # Predict on test data
# =====
y_pred_prob = model.predict(X_test)
y_pred = (y_pred_prob > 0.5).astype(int)

7/7 ————— 0s 3ms/step

```

```

In [97]: # Evaluate performance
# =====
test_acc = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {test_acc:.2f}")

```


Test Accuracy: 0.96

```

In [98]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("ANN Confusion Matrix")
plt.show()

```





```
In [99]: # Classification Report
print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0	0.93	0.98	0.96	102
1	0.98	0.93	0.96	103
accuracy			0.96	205
macro avg	0.96	0.96	0.96	205
weighted avg	0.96	0.96	0.96	205

```
In [100]: # Plot training history
# =====
plt.figure(figsize=(12,4))

# Accuracy
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Accuracy Over Epochs')
plt.legend()
```

```
# Loss
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Over Epochs')
plt.legend()

plt.show()
```



In [ ]:

In [ ]: